



IS322-Information Retrieval Course

Java Web Crawler with TF-IDF and Cosine Similarity

Name	ID
Israa Mohamed El-Sayed	20221252
Joseph Sameh Fouad	20220099
Youssef Joseph Adib	20220389
Jonathan Mokhles William	20220100
Amany Mohamed Hussein	20221026
Salma Mohamed Sarhan	20220152

Project overview:

Our project is a simple search engine for wikipedia pages

Main steps:

- 1)Crawl: visit web pages starting from two wikipedia links.
- 2)Build an inverted index: collect and organize the words from the pages.
- 3)User query: the user enters some search words.
- 4)Tf-idf and cosine similarity:we calculate how similar each page is to the query.
- 5)Ranking: show the top 10 most relevant pages based on similarity scores.

Web Crawler Module Explanation: by Israa Mohamed & Amany Mohamed

- Manages crawling of Wikipedia pages.
- Tracks:
 - ❖ visited set to avoid duplicate URLs.
 - ❖ docCounter to count the number of documents crawled.
- crawl() method:
 - ❖ Takes seedUrl (starting URL) and maxPages (maximum pages to crawl, set to 10).
 - ❖ Initializes a queue starting with the seed URL.
 - ❖ Skips URLs that:
 - Are already visited.
 - Do not start with "<https://en.wikipedia.org/wiki/>".
 - ❖ Fetches each valid URL using Jsoup.connect(currentUrl).get().
 - ❖ Extracts:
 - Page title using doc.title().
 - Paragraphs using doc.select("p").
 - ❖ Cleans text by:
 - Removing citations like [1], [note 3] using regex.
 - Skipping empty paragraphs.
 - Appending cleaned text to a StringBuilder.

Saving and Expanding the Crawl:

- Saves crawled content into a **Documents** folder:

- ❖ Each page saved as a separate text file (e.g., doc1.txt).
- ❖ File includes URL, page title, and cleaned paragraph text.
- Updates counters:
 - ❖ Increments docCounter and localSaved after each page.
- Extracts all hyperlinks (a:href) from the page.
- Adds only valid links to the queue:
 - ❖ Must belong to English Wikipedia domain.
 - ❖ Should contain important keywords (like Ancient, Akhenaten, etc.).
 - ❖ Must not already exist in the visited set or the queue.
- Continues crawling until 10 pages are successfully saved per seed.

Inverted Index Class Explanation: by Youssef Joseph & Joseph Sameh

Purpose of Inverted Index:

- Builds an inverted index from a collection of documents (text files or list of strings).
- Inverted Index maps each term (word) to a list of postings (document ID + term frequency).
- Used in search engines and information retrieval systems for fast keyword lookups.

Key Components:

- Map<String, List<Posting>> index -----> Stores the inverted index. Key = term, Value = list of postings
- Normalizer normalizer -----> Preprocesses the text: tokenization, lemmatization

Main Methods:

1- buildIndex(String folderPath)

- Reads all .txt files from a folder.
- Skips metadata lines (like "URL:" and "Title:").
- Normalizes and tokenizes the content.
- Builds term frequencies and updates the index.

2- buildIndex(List<String> documents)

- Builds the index directly from a list of in-memory document strings.
- Useful when documents are not in files.

3- Helper Methods

- **extractFileNumber(String name):** Extracts numbers from filenames for proper sorting.

- **printIndex()**: Prints the inverted index nicely formatted.
- **getDocumentCount()**: Returns how many unique documents are indexed.
- **getTerms()**: Returns the set of all terms.
- **getDocumentFrequency(String term)**: Returns how many documents contain a given term.
- **getTermFrequencies(int docId)**: Returns term frequencies for a specific document.
- **getAllDocumentIds()**: Returns the set of all document IDs.

How the Index is Built:

- 1-Read documents → ignore metadata.
- 2- Normalize text → lowercasing, cleaning, lemmatization.
- 3- Tokenize text into words.
- 4-Count term frequency (TF) for each document.
- 5-Update index: For each term, add or update its posting list.

TF-IDF Calculation & Ranking Explanation: by Jonathan Mokhles & Salma Mohamed

TF-IDF Calculator:

Fields

- **invertedIndex**: A map storing terms and their associated postings (document occurrences)
- **totalDocuments**: The total number of documents in the corpus
- **documentVectors**: Stores the final TF-IDF vectors for each term-document pair

Main Methods:

compute()

Calculates TF-IDF scores for all terms across all documents by:

1. Iterating through each term in the inverted index
2. Calculating IDF for the term
3. For each document containing the term:
 - Calculating term frequency (TF) weight
 - Computing TF-IDF score (TF × IDF)
 - Storing the result in documentVectors

```
calculateIDF(String term):
```

Computes Inverse Document Frequency using:

$IDF = \log_{10}(\text{totalDocuments} / \text{documentFrequency})$.

```
getTFIDF(String term, int docId):
```

Retrieves the precomputed TF-IDF score for a specific term-document pair.

```
printIndex():
```

Prints the TF-IDF index in a human-readable format.

Supporting Class: `DocumentTFIDF`:

A simple data structure that pairs:

- **docId:** Document identifier
- **tfidf:** Computed TF-IDF score for a term in this document

How Ranking Works:

1. Preprocessing:

- Query terms are processed (tokenized, maybe lemmatized).
- Document terms have already been processed and indexed with their TF-IDF scores.
- You fetch a list of documents and their tf-idf vectors for terms matching the query.

2. Build Document Vectors:

- You **group all tf-idf weights by document**.
- Each document is represented as a **vector**:
`{term → tf-idf weight}`

3. Normalize Document Vectors:

- Each document vector is normalized to unit length:
normalized weight = tf-idf weight / vector length.

- `normalizeVector(docVector)` is called for every document.
- Vector length = square root of the sum of squares of tf-idf weights.

4. Build Query Vector

- For the query:
 - For each **term**, get its **idf** (from corpus).
 - Compute **query weight** = $\text{tf} \times \text{idf}$ (usually $\text{tf} = 1$).
 - Where tf = frequency of the term in query.
- Then normalize the query vector **the same way** (divide by its vector length).

5. Compute Cosine Similarity:

- For each document:
 - Compute the **dot product** between the normalized **query vector** and the normalized **document vector**:

$$\text{cosine similarity} = \sum \text{terms in query} (\text{query weight} \times \text{document weight}).$$
 - for all terms in the query => compute query weight \times document weight.
- This gives a score for how similar the document is to the query.

6. Ranking:

- After computing scores for all documents:
 - Sort the documents by their **similarity scores in descending order**.
 - Return top-ranked documents as the **final ranking**.

Code Components Mapping:

Main Methods	Purpose
<code>Ranker.java/ buildNormalizedDocumentVectors()</code>	Creates and normalizes document vectors

<code>Ranker.java/ normalizeVector(Map<String, Double>)</code>	Divides each vector(query or document) component by the vector length
<code>CosineSimilarity.java/ calculate(queryVec, docVec)</code>	Computes cosine similarity via dot product
<code>Ranker/ rankAndDisplayTopDocuments() method</code>	Calls all the above, sorts the scores, returns ranked results

Simple Diagram:

Query Terms → build Query Vector → normalize



Document TFIDFs → build Doc Vectors → normalize



Compute dot products → Cosine Similarities



Sort by score → Ranked documents